

Technical Stack & Architectural Flow

Agentic AI Hiring Platform - Complete Technical Overview

TECHNICAL STACK

Backend Framework

- **FastAPI 0.109.0** - Modern, high-performance Python web framework
 - Async/await support for high concurrency
 - Automatic API documentation (Swagger UI & ReDoc)
 - Type validation with Pydantic
 - Fast request/response handling
- **Uvicorn 0.27.0** - ASGI server
 - Production-ready server
 - WebSocket support
 - HTTP/2 support

Database Layer

- **PostgreSQL** - Primary relational database
 - Robust ACID compliance
 - JSON support for complex data structures
 - Full-text search capabilities
- **SQLAlchemy 2.0.25** - ORM (Object-Relational Mapping)
 - Type-safe database operations
 - Connection pooling
 - Migration support
- **Alembic 1.13.1** - Database migrations
 - Version control for schemas
 - Automated migration scripts

AI & Machine Learning

- **HuggingFace Hub** - AI/ML model hosting
 - Access to transformer models
 - Inference API integration
 - Model: *sentence-transformers/all-MiniLM-L6-v2*
- **NumPy 1.26.3** - Numerical computing
 - Vector operations
 - Similarity calculations
 - Matrix operations

Document Processing

- **PyPDF2 3.0.1** - PDF text extraction
 - Resume parsing
 - Job description parsing
 - Multi-page document support

****Web & API****

- Python Multipart 0.0.6 - File upload handling

- Multipart form data processing
- Large file support

- Pydantic 2.5.3 - Data validation

- Type checking
- Schema validation
- Automatic documentation

- Python-dotenv 1.0.0 - Environment management

- Configuration management
- Secrets handling

****Deployment & Infrastructure****

- Docker - Containerization

- Consistent environments
- Easy deployment
- Scalability

- Render.com - Cloud hosting

- Auto-deployment from Git
- PostgreSQL hosting
- Environment variables

****Security & Performance****

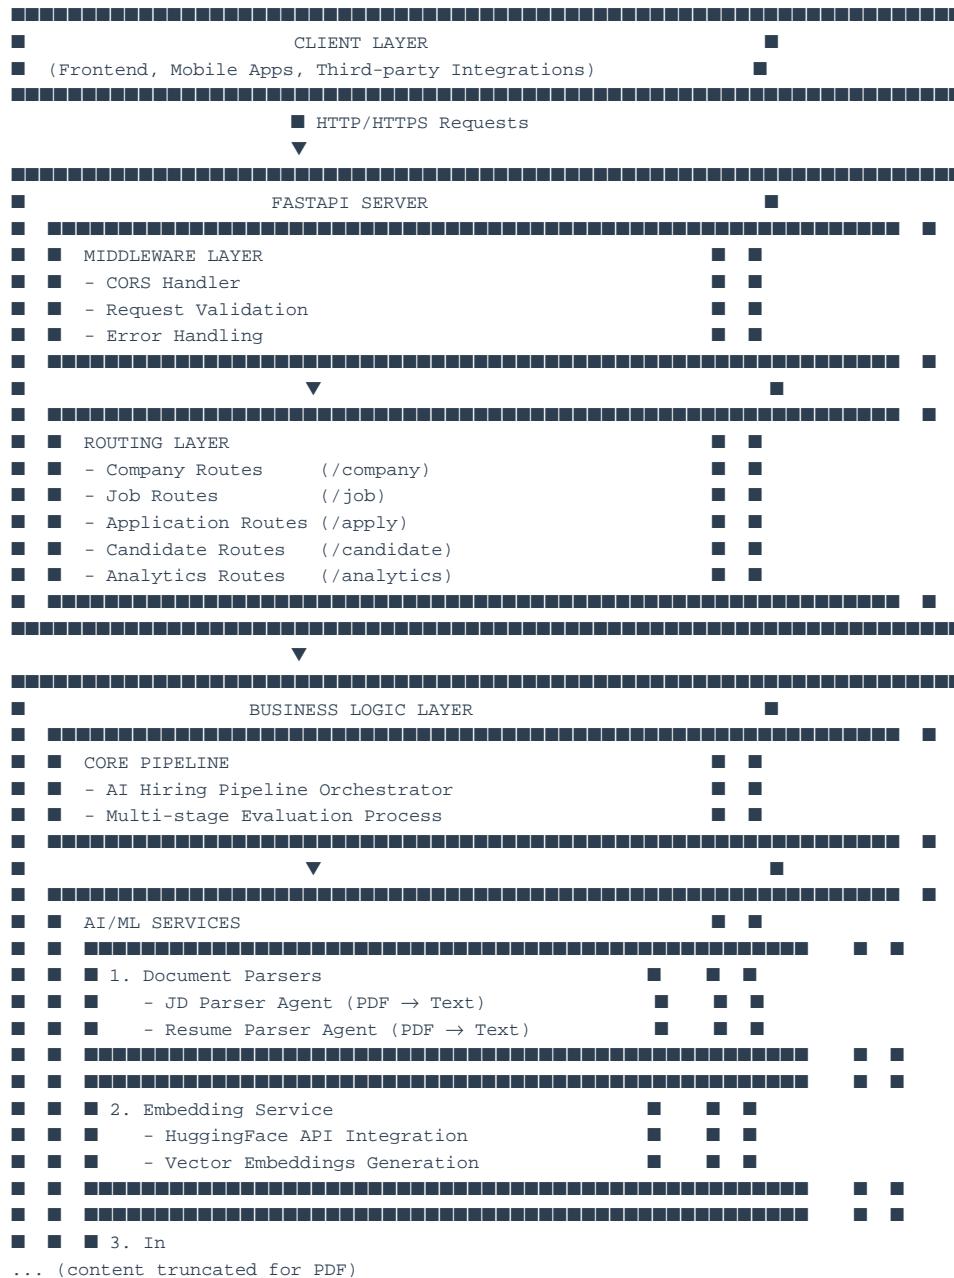
- CORS Middleware - Cross-Origin Resource Sharing

- Connection Pooling - Database optimization

- Pre-ping Health Checks - Connection reliability

ARCHITECTURAL FLOW

System Architecture Overview



DETAILED PROCESS FLOWS

****1. Job Posting Flow****

```
graph TD; A[Recruiter Uploads JD PDF] --> B[JD Parser Agent Extracts Text]; B --> C[Text Cleaning & Normalization]; C --> D[Embedding Service Generates Vector]; D --> E[Inference Engine Extracts Skills]; E --> F[Skills Classified (Technical/Soft)]; F --> G[Required vs Nice-to-Have Detection]; G --> H[Job Record Created in Database]; H --> I[Job ID Returned to Client]
```

The diagram illustrates a sequential process for job search. It starts with a Recruiter uploading a JD PDF, which is then processed by a JD Parser Agent to extract text. This text undergoes Text Cleaning & Normalization. An Embedding Service generates a vector from the cleaned text. An Inference Engine extracts skills from the vector. The skills are then classified as Technical or Soft. Finally, a Required vs Nice-to-Have Detection step leads to a Job Record being created in a database, and a Job ID is returned to the client.

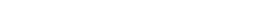
Key Technologies:

- PyPDF2 for PDF parsing
 - Text cleaning utilities
 - HuggingFace embeddings
 - AI skill extraction
 - PostgreSQL storage

****2. Application Submission & Evaluation Flow****

- RFS: Cosine Similarity (JD vs Resume embeddings)
 - DCS: Weighted Skill Match (Technical skills only)
 - ELC: Experience Level Compatibility
 - Composite: Weighted Average
 - ↓
 - [3.2] Fraud Detection Analysis
 - Compare with existing candidates
 - Resume similarity check (>90% threshold)
 - Email duplication check
 - Generate fraud report
 - ↓
 - [3.3] Decision Service
 - Evaluate composite score against thresholds
 - Consider fraud flags
 - Generate decision (Selected/Rejected/Review)
 - Create decision reasoning
 - ↓
 - [3.4] Explanation Generation
 - Identify strengths & weaknesses
 - Provide actionable recommendations
 - Generate detailed explanation
 - ↓
 - [3.5] XAI Analysis (for transparency)
 - Break down decision factors
 - Assign contribution percentages
 - Generate confidence levels
 - ↓
 - [3.6] Skill Gap Analysis
 - Identify matched vs missing skills
 - Separate required from nice-to-have gaps
 - Generate learning roadmap
 - Estimate closure time
 - Identify transferable skills
 - ↓
 - [3.7] Skill Evidence Graph Generation
 - Create node-edge graph structure
 - Visual representation of skill relationships
 - ↓

■ STAGE 4: RANKING & STORAGE



↓
Store Application with All Metadata

↓
Log Audit Trail

■ STAGE 5: RESPONSE GENERATION

- Application ID
- All Scores (RFS, DCS, ELC, Composite)
- Ranking Position
- Decision & Reasoning
- Explanation
- Fraud Analysis

Processing Time: ~2-5 seconds per application

****3. Candidate Master Endpoint Flow****

Request: GET /candidate/{id}/master

1

Database Query: Fetch Candidate

```
↓  
Database Query: Fetch All Applications  
↓  
For Each Application:  
  ■■ Fetch Job Details  
  ■■ Fetch Company Details  
  ■■ Fetch All Scores  
  ■■ Fetch Decision & Reasoning  
  ■■ Fetch Fraud Analysis  
  ■■ Fetch Skill Match Data  
  ↓  
Calculate Summary Statistics  
  ■■ Total Applications  
  ■■ Selected/Rejected/Pending Counts  
  ■■ Average Composite Score  
  ■■ Best Application  
  ↓  
Compile Comprehensive Response  
  ↓  
Return JSON with Complete Data
```

****4. Analytics & Reporting Flow****

```
graph TD; A[Request Analytics Endpoint] --> B[XAI Explanation]; A --> C[Skill Gap Analysis]; A --> D[Company Dashboard]; B --> B1[Fetch application data]; B --> B2[Calculate factor contributions]; B --> B3[Generate confidence levels]; B --> B4[Return transparent explanation]; C --> C1[Identify missing skills]; C --> C2[Classify by priority]; C --> C3[Generate learning roadmap]; C --> C4[Calculate closure time]; D --> D1[Aggregate all company jobs]; D --> D2[Calculate hiring metrics]; D --> D3[Generate trends analysis]; D --> D4[Return comprehensive dashboard]
```

The diagram illustrates a process flow starting from a central point:

- Request Analytics Endpoint** (represented by a downward arrow) leads to three parallel paths:
- XAI Explanation**:
 - Fetch application data
 - Calculate factor contributions
 - Generate confidence levels
 - Return transparent explanation
- Skill Gap Analysis**:
 - Identify missing skills
 - Classify by priority
 - Generate learning roadmap
 - Calculate closure time
- Company Dashboard**:
 - Aggregate all company jobs
 - Calculate hiring metrics
 - Generate trends analysis
 - Return comprehensive dashboard

SCORING ALGORITHMS

1. Role Fit Score (RFS)

```
RFS = cosine_similarity(JD_embedding, Resume_embedding)
```

Where:

- JD_embedding = 384-dimensional vector from sentence-transformers
- Resume_embedding = 384-dimensional vector from sentence-transformers
- Result range: 0.0 to 1.0 (0% to 100%)

2. Domain Competency Score (DCS)

WEIGHTED SKILL MATCHING:

Step 1: Classify JD skills

- Required skills (must-have)
- Nice-to-have skills (optional)

Step 2: Match candidate skills

- Direct matches
- Synonym matches
- Related skill matches

Step 3: Calculate weighted score

```
DCS = (Required_Match_Score × 0.7) + (NiceToHave_Match_Score × 0.3)
```

Where:

```
Required_Match_Score = (Matched Required / Total Required) × 100
```

```
NiceToHave_Match_Score = (Matched NiceToHave / Total NiceToHave) × 100
```

Note: Only technical skills counted, soft skills excluded

3. Experience Level Compatibility (ELC)

```
If candidate_exp >= required_exp:  
    ELC = 1.0 (100% match)  
Else if candidate_exp >= required_exp × 0.75:  
    ELC = 0.8 (80% match)  
Else if candidate_exp >= required_exp × 0.5:  
    ELC = 0.6 (60% match)  
Else:  
    ELC = 0.3 (30% match - significant gap)
```

Additional factors:

- Relevant experience bonus (+10%)
- Over-qualification penalty (-5%)

4. Composite Score

```
Composite = (RFS × 0.25) + (DCS × 0.50) + (ELC × 0.25)
```

Weights explanation:

- 25% for semantic fit (RFS)
- 50% for skill matching (DCS) - most important
- 25% for experience (ELC)

****5. Decision Thresholds****

```
If fraud_flag == True:  
    Decision = "Rejected - Fraud Detected"  
Else if composite >= 0.75:  
    Decision = "Selected"  
Else if composite >= 0.60:  
    Decision = "Review Required"  
Else:  
    Decision = "Rejected"
```

SECURITY FEATURES

1. Environment Variables

- Sensitive data in .env files
- No hardcoded credentials
- DATABASE_URL encryption

2. Database Security

- Connection pooling with pre-ping
- SQL injection prevention (ORM)
- ACID transactions

3. API Security

- CORS configuration
- Request validation
- Error handling & logging

4. Fraud Detection

- Resume similarity checks
- Duplicate email detection
- Suspicious pattern identification

DATA MODELS

Company

```
- id: Integer (PK)
- name: String
- description: Text
- created_at: DateTime
```

Job

```
- id: Integer (PK)
- company_id: Integer (FK)
- role: String
- location: String
- salary: String
- employment_type: String
- required_experience: Integer
- jd_text: Text
- jd_embedding: JSONB (384-dim vector)
- skills_extracted: JSONB
- created_at: DateTime
```

Candidate

```
- id: Integer (PK)
- name: String
- email: String (indexed)
- mobile: String
- linkedin: String
- github: String
- experience: Integer
- resume_text: Text
- resume_embedding: JSONB (384-dim vector)
- skills_extracted: JSONB
- created_at: DateTime
```

Application

```
- id: Integer (PK)
- job_id: Integer (FK)
- candidate_id: Integer (FK)
- rfs: Float (Role Fit Score)
- dcs: Float (Domain Competency Score)
- elc: Float (Experience Level Compatibility)
- composite_score: Float (indexed)
- rank: Integer (indexed)
- similarity_index: Float
- fraud_flag: Boolean (indexed)
- fraud_details: JSONB
- decision: String (indexed)
- decision_reason: Text
- explanation: JSONB
- skill_match: JSONB
- experience_details: JSONB
- status: String (indexed)
```

```
- created_at: DateTime (indexed)
```

****AuditLog****

```
- id: Integer (PK)
- entity_type: String (Job/Candidate/Application)
- entity_id: Integer
- action: String
- details: JSONB
- timestamp: DateTime
```

DEPLOYMENT ARCHITECTURE

Docker Container

- Base: Python 3.11-slim
- Port: 10000
- User: Non-root (appuser)
- Dependencies: Installed from requirements.txt

Render.com Configuration

- Auto-deploy from GitHub
- PostgreSQL database instance
- Environment variables management
- Health checks enabled
- Auto-scaling support

Environment Variables Required

```
DATABASE_URL=postgresql://...  
HF_API_KEY=hf_...  
SIMILARITY_THRESHOLD=0.90  
ENVIRONMENT=production  
LOG_LEVEL=INFO
```

PERFORMANCE CHARACTERISTICS

- **API Response Time:** 50-200ms (without AI processing)
- **Application Evaluation:** 2-5 seconds
- **Concurrent Requests:** 100+ (with proper scaling)
- **Database Queries:** Optimized with indexes
- **Embedding Generation:** ~500ms per document
- **Skill Extraction:** ~1-2 seconds per document

FUTURE ENHANCEMENTS (Optional Dependencies)

- **Redis**: Caching frequently accessed data
- **Celery**: Background task processing
- **JWT Authentication**: User authentication
- **Rate Limiting**: API request throttling
- **Sentry**: Error tracking & monitoring

API DOCUMENTATION

- **Swagger UI:** [`/docs`](#) endpoint
- **ReDoc:** [`/redoc`](#) endpoint
- **OpenAPI Schema:** Auto-generated
- **PDF Documentation:** Available in repository

This architecture provides a robust, scalable, and AI-powered hiring platform with transparent decision-making and comprehensive candidate evaluation.